

Purpose: This program will explore the utilization of the buffered I/O algorithm to efficiently read/write large quantities of data.

Once working, this program will then be used to compare I/O operation times using different size buffers (See Part B).

The program must be written from scratch.

Program Specifications:

Write an assembly program that will read in a file and outputs a file with the hex values of the first file.

The program will accept filenames as input from the command line arguments. Additionally, the program can include an optional "-p" command line argument that indicates the program should print each byte read in.

Example Command Lines:

```
./a.out inputFile.dat outputFile.txt
```

```
./a.out inputFile.dat outputFile.txt -p
```

The -p option can only occur in the last command line argument.

The program should verify the command line count (print a usage guide if there is only a single command line argument).

Open the input and output files and verify that they opened successfully. Otherwise, output an appropriate error message and end the program.

The program should then repeatedly call `getBytes` and `writeBytes` until the input file runs out of data. For each byte of data read in, the write function should expand it to five characters: "0xHH", `LINEFEED` where the two H are replaced by the hex value for that byte. This means the output file will be 5 times larger than the input. The output buffer should be 5 times larger as well. Once the

end of the input file is reached, the program should use the `finalizeOutputBuffer` function to write any remaining data in the output buffer to the output file.

This program will require the use of two buffers, one for input and one for output. The output buffer should be five times as large as the input.

If the user uses the `-p` command line option, the program should additionally call `printByte` using each byte read in. This will severely affect the performance of the program as console I/O operations are slow.

Use buffer sizes of 100,000 (input) and 500,000 (output).

Once completed, use the program to answer part B. Make sure to set the buffer sizes back to 100,000 and 500,000 before turning in your program.

Required Functions

getBytes

Argument 1: Address to file descriptor

Argument 2: Address to input buffer

Argument 3: 64 bit integer value for Buffer Size

Argument 4: Address to buffered bytes count

Argument 5: Address to read bytes count

Argument 6: Address to a byte value for tracking endOfFile

Argument 7: Address to string file access error message

Argument 8: Address to store the byte read in from the buffer

Returns -1 if the operation failed, 0 if the file has no more data, or 1 if the operation succeeded

Reads in one byte using the buffered I/O algorithm.

writeByte

Argument 1: Address to file descriptor

Argument 2: Address to output buffer

Argument 3: 64 bit integer value for Buffer Size

Argument 4: Address to buffered bytes count

Argument 5: Byte value to expand and write to file

Argument 6: Address to string file access error message

Returns -1 if the operation failed, or 1 if the operation succeeded

Converts and writes five characters using the buffered I/O algorithm. When the buffer is full, the function should write it to the file and clear the buffer.

Dividing the byte (argument 5) by 16 will give you the two hex values required.

Examples

$$124 / 16 = 7 \text{ r } 12$$

$$7 / 16 = 0 \text{ r } 7$$

0x7C

$$255 / 16 = 15 \text{ r } 15$$

0xFF

Each hex value will need to be adjusted to its appropriate ASCII representation. Check to see if the value is less than 10 or greater than 10 to determine which offset value to use.

7

12

$$7 < 10$$

$$7 + '0' = '7'$$

$$12 > 10$$

$$12 + 'A' - 10 = 'C'$$

Each byte should be converted into 5 characters:
"0x", the two converted hex digits, linefeed

Once the buffer is full, attempt to write it to the given file and reset the buffer (buffered characters = 0).

finalizeOutputBuffer

Argument 1: Address to file descriptor

Argument 2: Address to output buffer

Argument 3: Address to buffered bytes count

Returns -1 if the operation failed, or 1 if the operation succeeded

Writes any remaining characters in the output buffer to the file.

printByte

Argument 1: an unsigned integer byte by value

Prints the byte in hex format to the console. The string must be stored as a local variable to printByte.

Part B:

In a document file (.docx or .odt), experiment with the buffer size by running the program with buffer sizes of 1, 1000, and 100,000. For each buffer size, run the program using the Linux time program on the combined.txt file. It may take several minutes to run.

```
time a.out combined.txt combinedOutput.txt
```

Create a table using the real, user, and system times and calculate the percentage of time each buffer size spends in system time(sys).

```
(sys time)/(real time)*100%
```

Explain why certain buffer sizes finish faster than others in 200 words or less (preferably much less).

Assembly and Linkage

Make sure to use main: instead of _start: as your starting function and use **g++ -g -no-pie assemblyFile.o** to link your program.

Submission

Once you are satisfied with the program, upload the assembly source code (.asm) file to the class website to Assignment #7 Part A and the document file to Assignment #7 Part B.

Example Executions

```
$ ./a.out inputFile.dat outputFile.txt  
$
```

```
$ ./a.out test.txt testOutput.txt -p  
0x54  
0x65  
0x73  
0x74  
0x69  
0x6E  
0x67  
0x2C  
0x20  
0x74  
0x65  
0x73  
0x74  
0x69  
0x6E  
0x67  
0x2E  
0x0A  
0x48  
0x65  
0x6C  
0x6C  
0x6F  
0x2C  
0x20  
0x57  
0x6F  
0x72  
0x6C  
0x64  
0x21  
$
```
